

Online Meetup

# Multi-Table Transactions with LakeFS and Delta Lake



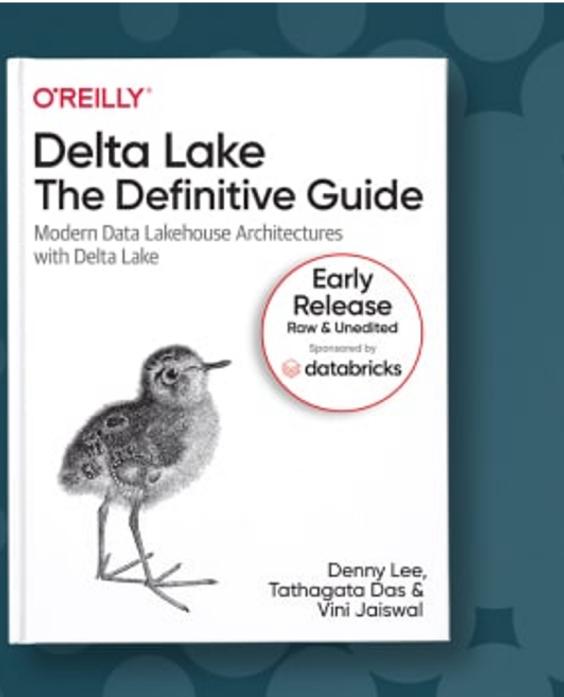
# About the Speaker



Denny Lee is a Developer Advocate at Databricks. He is a hands-on distributed systems and data sciences engineer with extensive experience developing internet-scale infrastructure, data platforms, and predictive analytics systems for both on-premise and cloud environments.

# O'Reilly Delta Lake: The Definitive Guide (Early Release)

Early Release  
**O'Reilly**  
**Definitive**  
**Guide**



<https://dbricks.co/dldg>

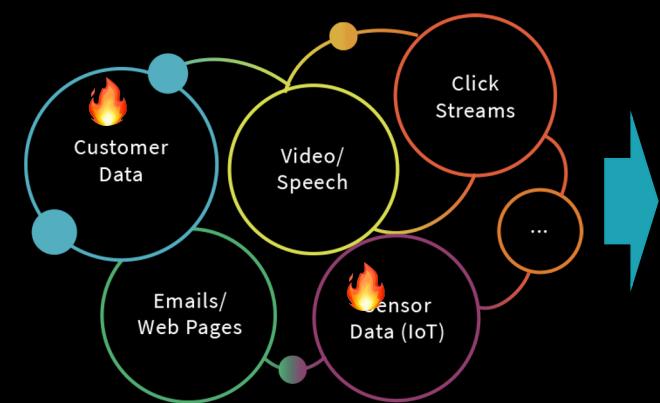
# About the Speaker



Paul Singman is a Developer Advocate at lakeFS. He previously led a Platform Engineering team at Equinox Fitness, where he worked on operational and user-facing analytic services. When not staring at screens, he likes to run, golf, play ping-pong and backgammon.

# The Promise of the Data Lake

## 1. Collect Everything



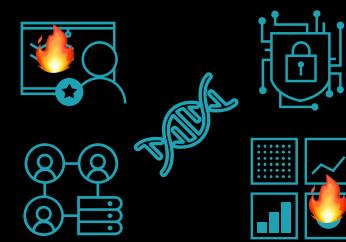
Garbage In

## 2. Store it all in the Data Lake



Garbage Stored

## 3. Data Science & Machine Learning



- Recommendation Engines
- Risk, Fraud Detection
- IoT & Predictive Maintenance
- Genomics & DNA Sequencing

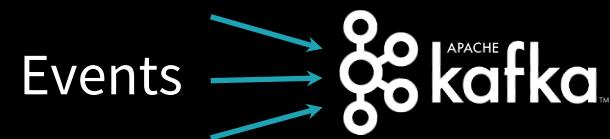
Garbage Out



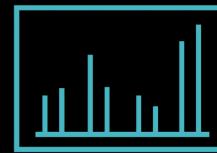
What does a typical  
**data lake** project look like?



# Evolution of a Cutting-Edge Data Lake



Data Lake



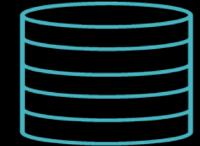
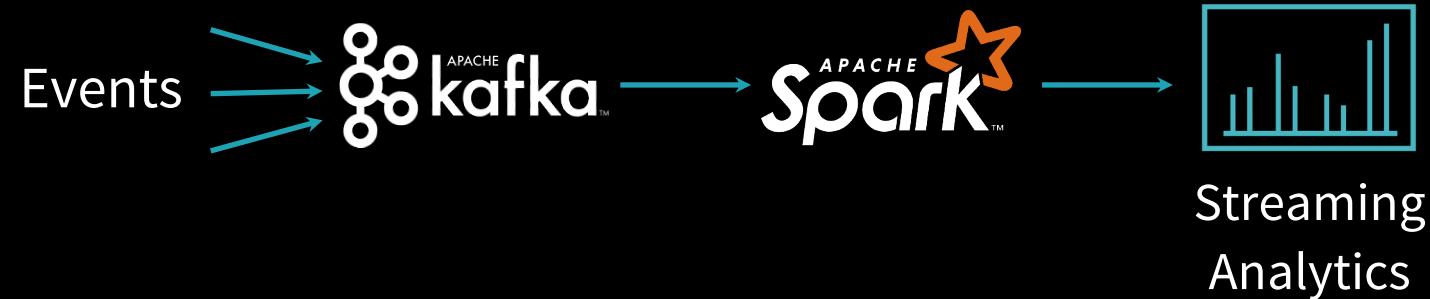
Streaming  
Analytics



AI & Reporting



# Evolution of a Cutting-Edge Data Lake



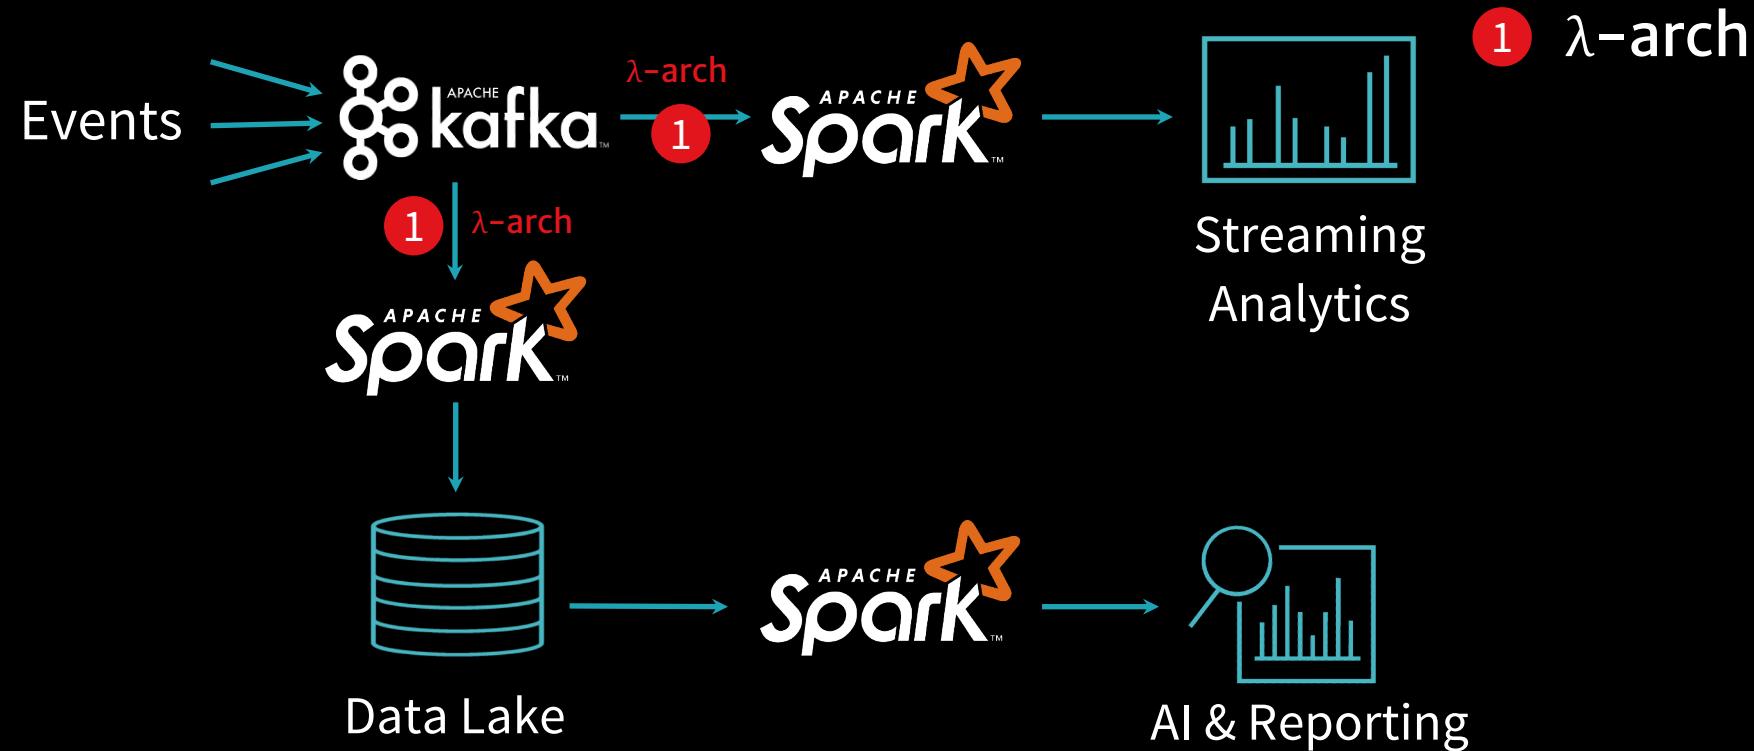
Data Lake



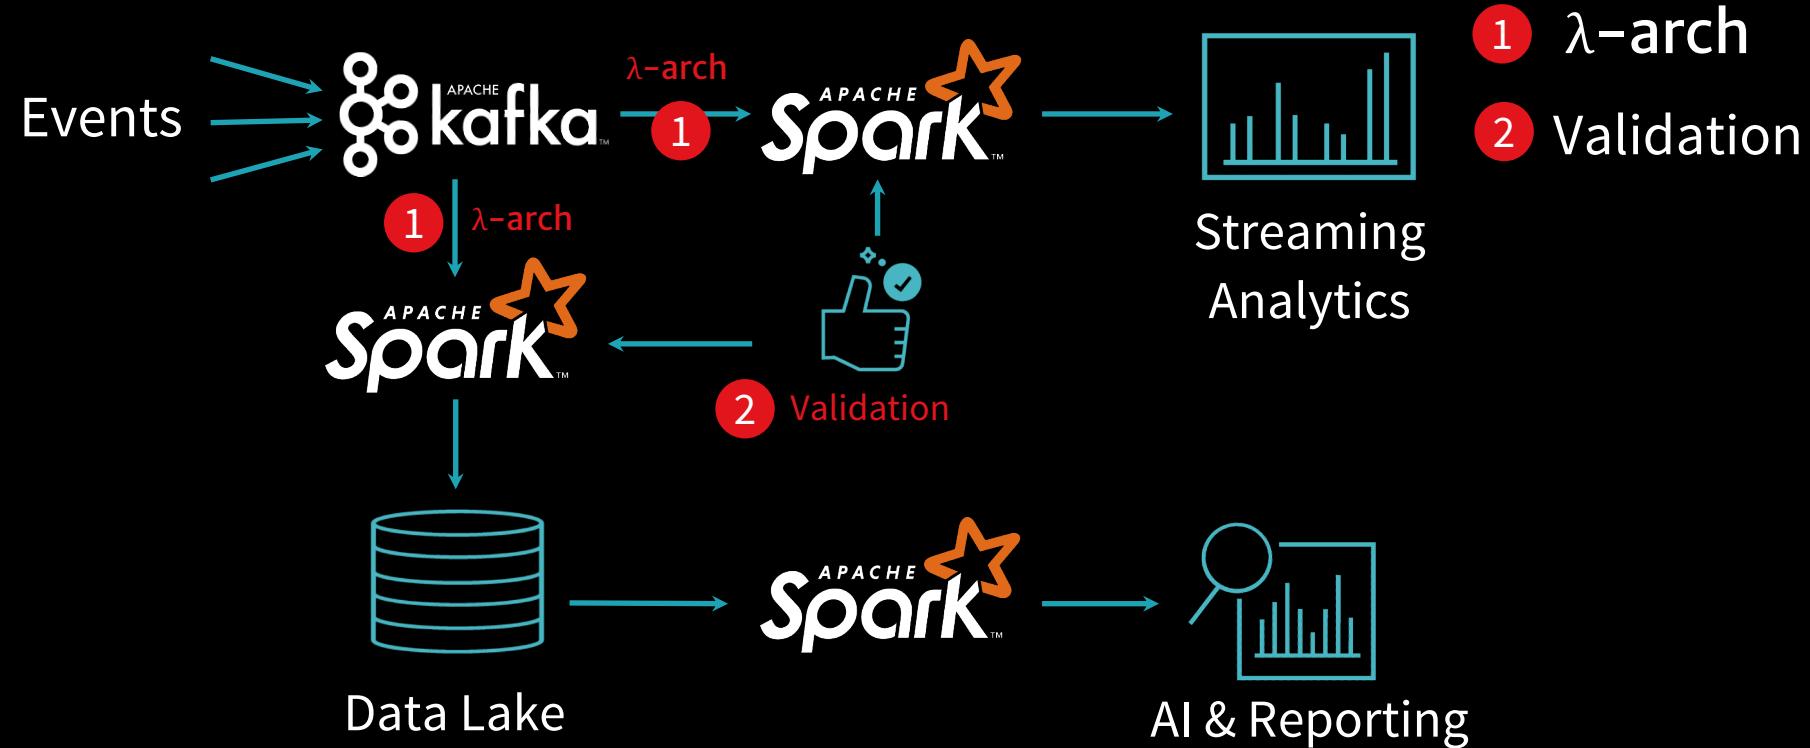
AI & Reporting



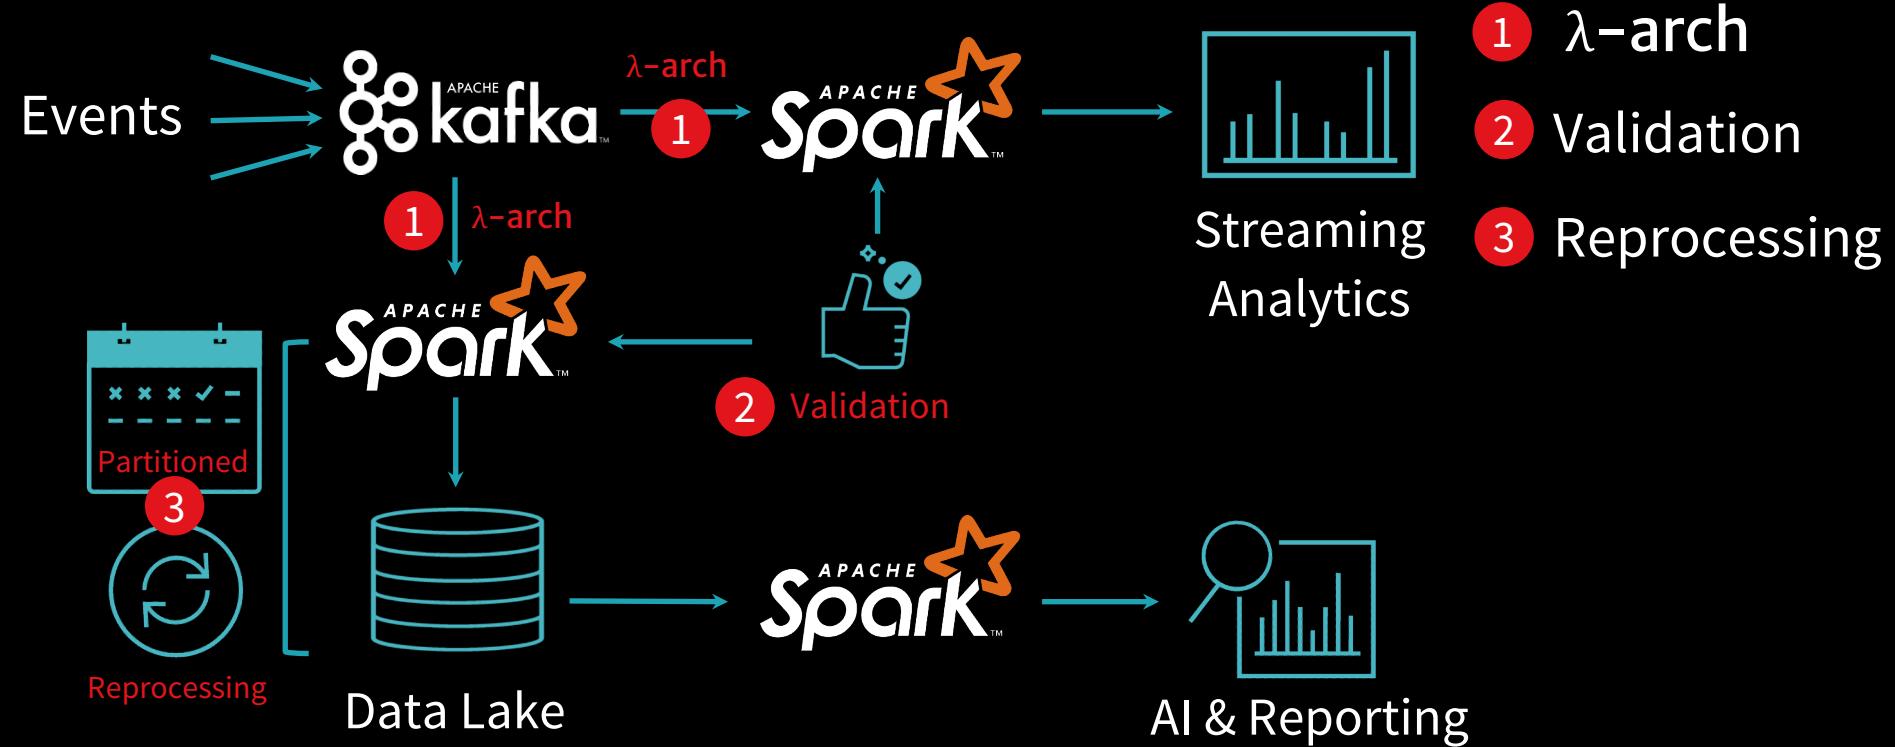
# Challenge #1: Historical Queries?



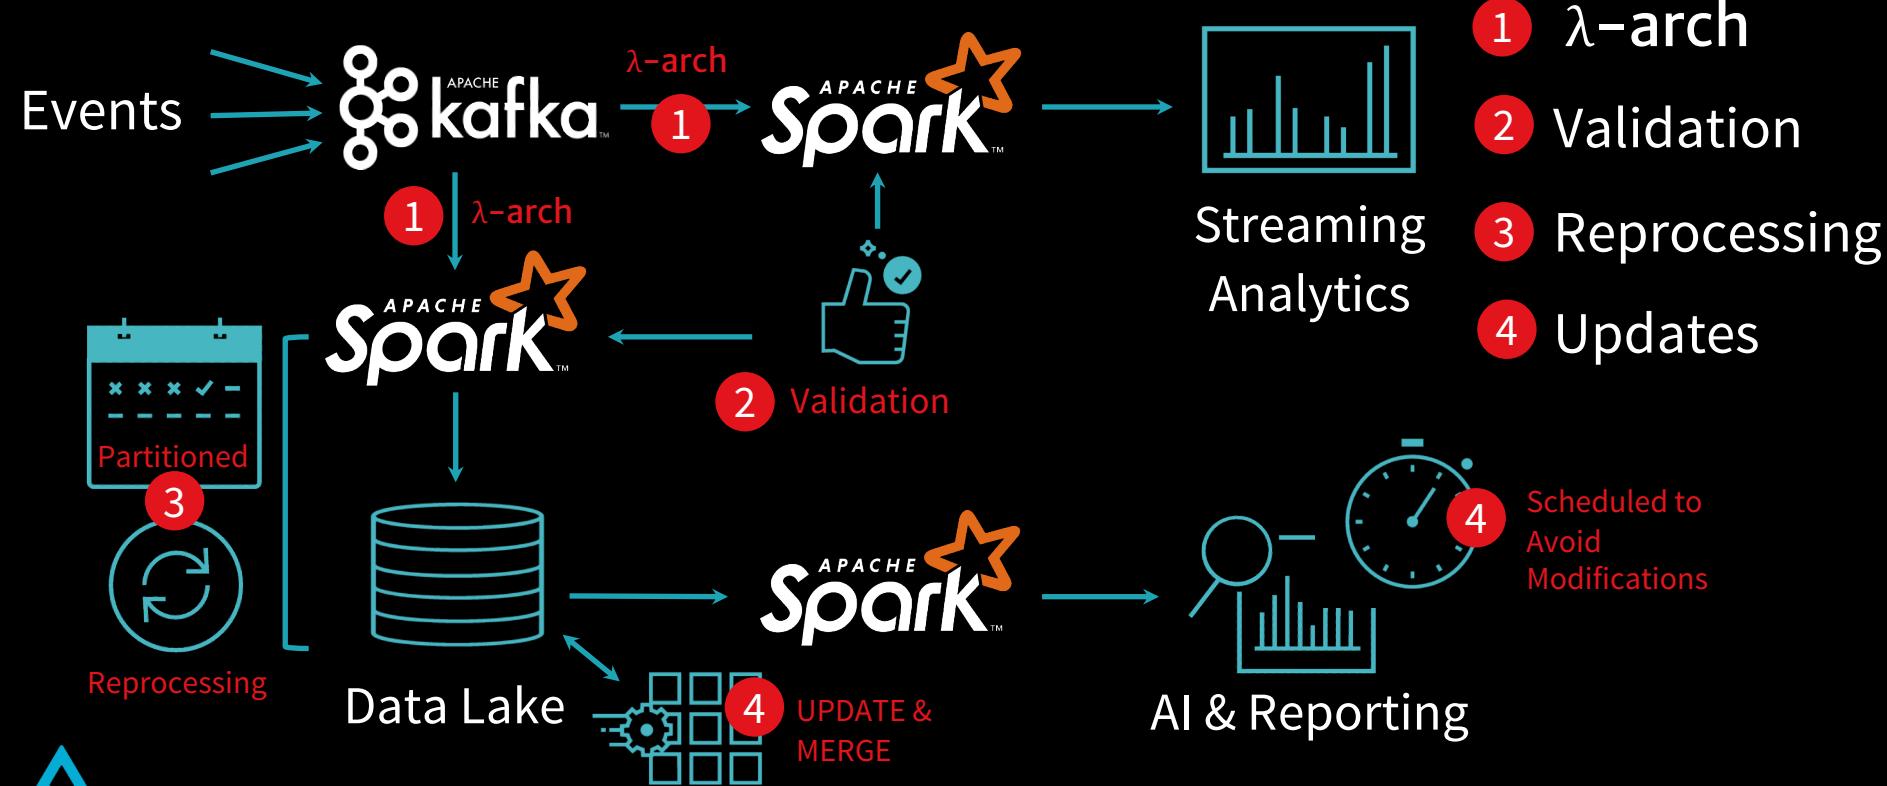
# Challenge #2: Messy Data?



# Challenge #3: Mistakes and Failures?



# Challenge #4: Updates?



Wasting Time & Money

Solving Systems Problems

Instead of Extracting Value From Data



# Data Lake Distractions



**No atomicity** means failed production jobs leave data in corrupt state requiring tedious recovery



**No quality enforcement** creates inconsistent and unusable data



**No consistency / isolation** makes it almost impossible to mix appends and reads, batch and streaming



# OSS Delta Lake Key Features

## Feature

---

ACID Transactions	Delta Lake brings ACID transactions to your data lakes. It provides serializability, the strongest level of isolation level. Learn more at <a href="#">Diving into Delta Lake: Unpacking the Transaction Log</a> .
Scalable Metadata Handling	Delta Lake can handle petabyte-scale tables with billions of partitions and files at ease.
Time Travel (data versioning)	Delta Lake provides data snapshots to <a href="#">access and revert to earlier versions</a> of data for audits, rollbacks or to reproduce experiments.
Open Format	All data in Delta Lake is stored in Apache Parquet format enabling Delta Lake to leverage the efficient compression and encoding schemes that are native to Parquet
Unified Batch and Streaming Source and Sink	A table in Delta Lake is both a batch table, as well as a streaming source and sink. Streaming data ingest, batch historic backfill, and interactive queries all just work out of the box.



# OSS Delta Lake Key Features (Continued)

## Feature

---

### Schema Enforcement and Evolution

Delta Lake provides the ability to specify your schema and enforce it. This helps ensure that the data types are correct and required columns are present, preventing bad data from causing data corruption. For more information, refer to [Diving Into Delta Lake: Schema Enforcement & Evolution](#).

### Audit History

Delta Lake transaction log records details about every change made to data providing a full audit trail of the changes.

### DML Operations

Delta Lake supports SQL, [Scala / Java](#) and [Python](#) APIs to merge, update and delete datasets allowing you to easily comply with GDPR and CCPA and simplifying use cases like change data capture. For more information, refer to [Diving Into Delta Lake: DML Internals](#)

### 100% Compatible with Apache Spark API

Developers can use Delta Lake with their existing data pipelines with minimal change as it is fully compatible with Spark, the commonly used big data processing engine.

---



# Delta Lake CY21H2 Roadmap

<https://github.com/delta-io/delta/issues/748>



- Delta Standalone Reader (Q3)
- Delta Standalone Writer (Q4)



- Flink/Delta Source (Q4)
- Flink/Delta Sync (22Q1)



- Pulsar/Delta Source (Q4)
- Pulsar/Delta Sync (22Q1)



- Spark 3.2 Support (Q4)



Hive3 Connector (Q3)



LakeFS Integration (Q3)

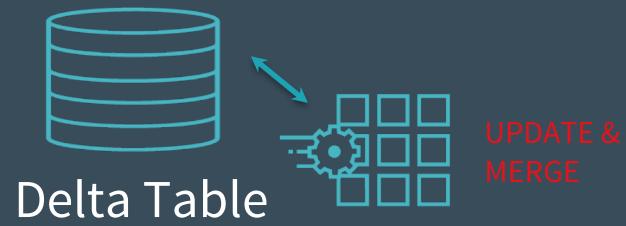


Nessie Integration (Q4)

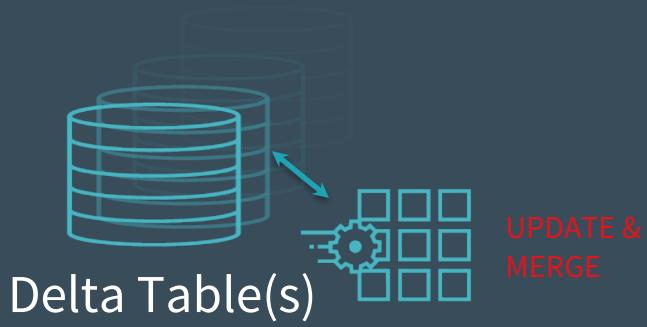


**Community meetings:** every two weeks have meetings with Trino, Flink, Nessie, LakeFS, Pulsar, Rust, and Core

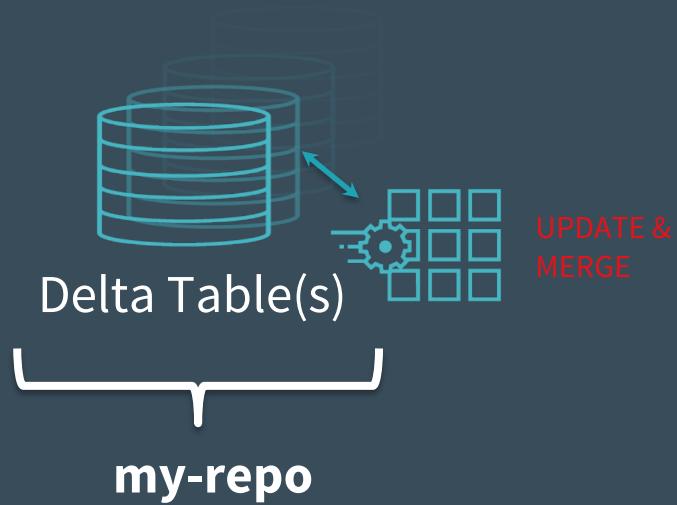
# Single-Table Consistency



# Multi - Table Consistency

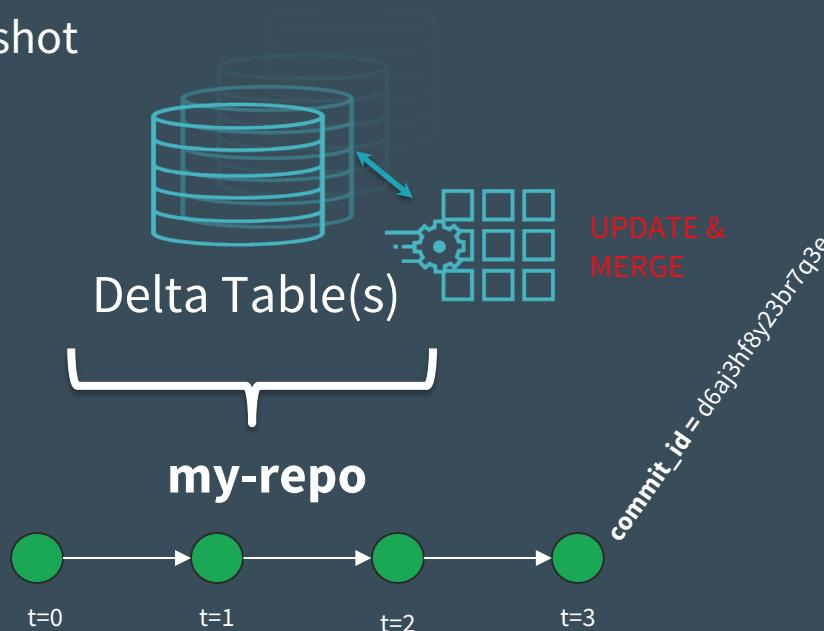


# Multi - Table Consistency



# Multi - Table Consistency

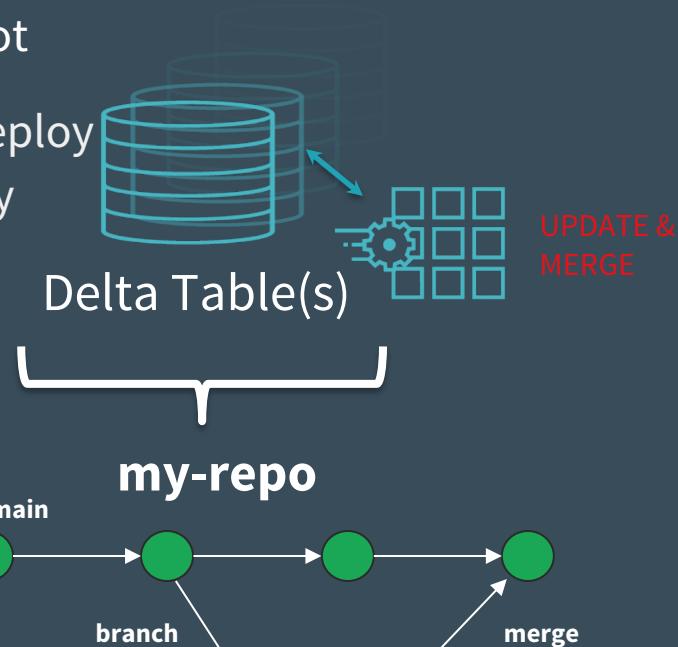
1. **commit** - save cross collection snapshot



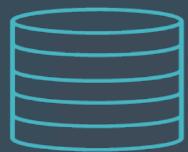
# Multi - Table Consistency

1. **commit** - save cross collection snapshot

2. **branch/merge** - deploy data atomically



# How it works – Metadata All the Way Down



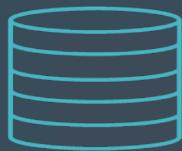
Transaction Log  
Table Versions

(Optional) Partition Directories  
Data Files

my\_table/  
  → \_delta\_log/  
    → 00000.json  
    → 00001.json  
  → date=2019-01-01/  
    → file-1.parquet



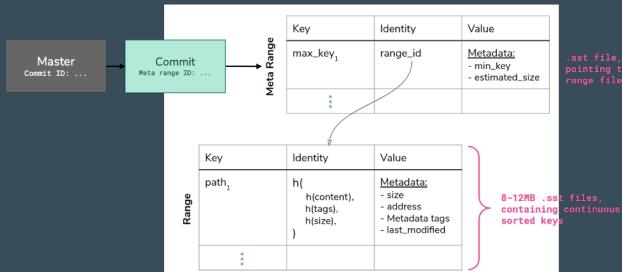
# How it works – Metadata All the Way Down



Transaction Log  
Table Versions

(Optional) Partition Directories  
Data Files

my\_table/  
  →\_delta\_log/  
  →0000.json  
  →0001.json  
  date=2019-01-01/  
  ↓file-1.parquet



Object Store



aws  
Google Cloud



MINIO

# Configuring lakeFS + Delta



# Configuring lakeFS + Delta

## 1. Set Spark configs to S3 API endpoint

```
spark.hadoop.fs.s3a.bucket.<repo-name>.access.key AKIAIOSFODNN7EXAMPLE  
spark.hadoop.fs.s3a.bucket.<repo-name>.secret.key wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY  
spark.hadoop.fs.s3a.bucket.<repo-name>.endpoint https://lakefs.example.com  
spark.hadoop.fs.s3a.path.style.access true
```

# Configuring lakeFS + Delta

## 1. Set Spark configs to S3 API endpoint

```
spark.hadoop.fs.s3a.bucket.<repo-name>.access.key AKIAIOSFODNN7EXAMPLE  
spark.hadoop.fs.s3a.bucket.<repo-name>.secret.key wJalrXUtnFEMI/K7MDENG/bPxRfiCYEXAMPLEKEY  
spark.hadoop.fs.s3a.bucket.<repo-name>.endpoint https://lakefs.example.com  
spark.hadoop.fs.s3a.path.style.access true
```

## 2. Update object store location paths

```
Cmd 8  
  
1 %sql  
2 DROP TABLE IF EXISTS d_loan_states;  
3  
4 CREATE TABLE d_loan_states  
5 USING delta  
6 LOCATION 's3a://bi-reports/dev-reports/tables/loan_by_state'  
7 AS SELECT * FROM loan_by_state;
```



The diagram shows two labels, "repo" and "branch", positioned above the "LOCATION" and "AS SELECT" clauses respectively. Two arrows originate from these labels and point towards their respective corresponding clauses in the SQL code.

▶ (3) Spark Jobs

# Recap

- Add your Delta tables to a lakeFS repository.
- Utilize git-inspired operations to provide atomicity on transactions across multiple tables.

Who is using  **DELTA LAKE**?



# Used by 5000+ of organizations world wide

> 1+ exabyte processed every three days



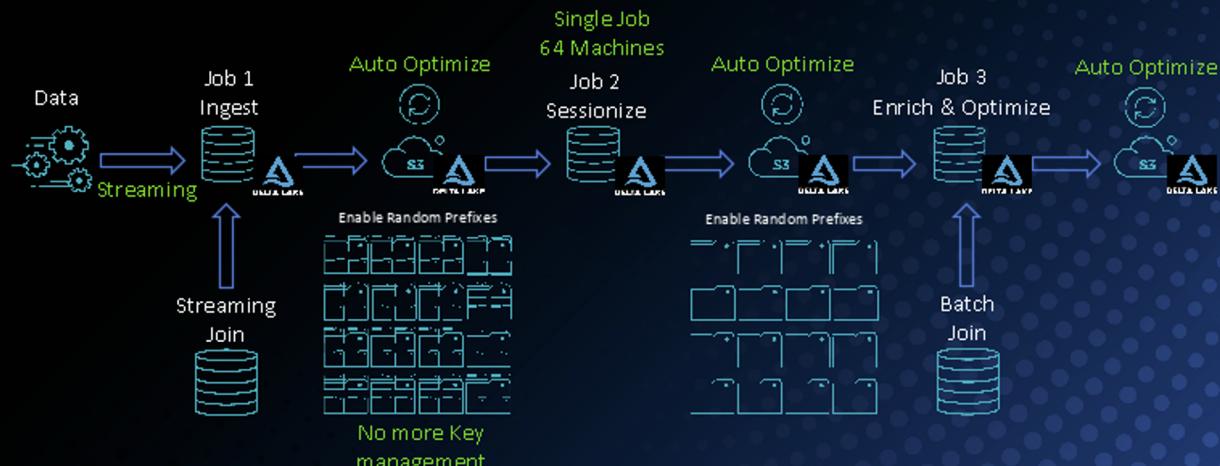
Barracuda





# COMCAST

## SESSIONIZATION WITH DELTA LAKE



FASTER QUERIES, RELIABLE PIPELINES, 10X REDUCTION IN COMPUTE!



Improved reliability:  
Petabyte-scale jobs

10x lower compute:  
640 instances to 64!

Simpler, faster ETL:  
84 jobs → 3 jobs  
halved data latency



How do I use  **DELTA LAKE** ?



# Get Started with Delta using Spark APIs

## Add Spark Package

```
pyspark --packages io.delta:delta-core_2.12:0.1.0  
bin/spark-shell --packages io.delta:delta-core_2.12:0.1.0
```

## Maven

```
<dependency>  
  <groupId>io.delta</groupId>  
  <artifactId>delta-core_2.12</artifactId>  
  <version>0.1.0</version>  
</dependency>
```

Instead of **parquet**...

```
dataframe  
.write  
.format("parquet")  
.save("/data")
```

... simply say **delta**

```
dataframe  
.write  
.format("delta")  
.save("/data")
```

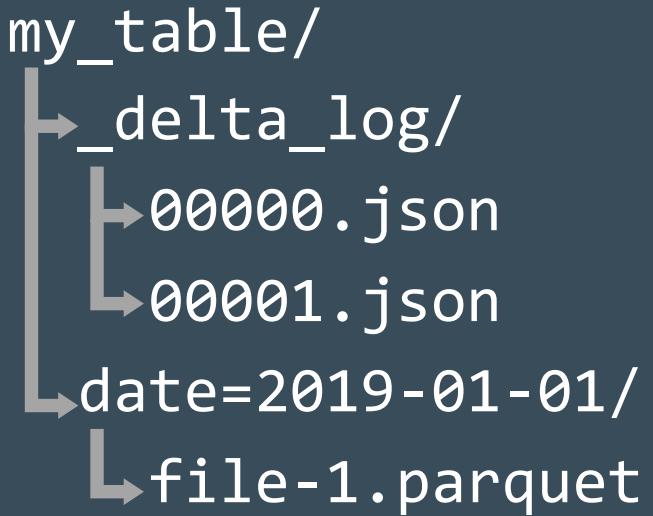


# How does DELTA LAKE work?



# Delta On Disk

Transaction Log  
Table Versions  
(Optional) Partition Directories  
Data Files



# Version N = Version N-1 + Actions

Change Metadata – name, schema, partitioning, etc

Add File – adds a file (with optional statistics)

Remove File – removes a file

**Result:** Current Metadata, List of Files



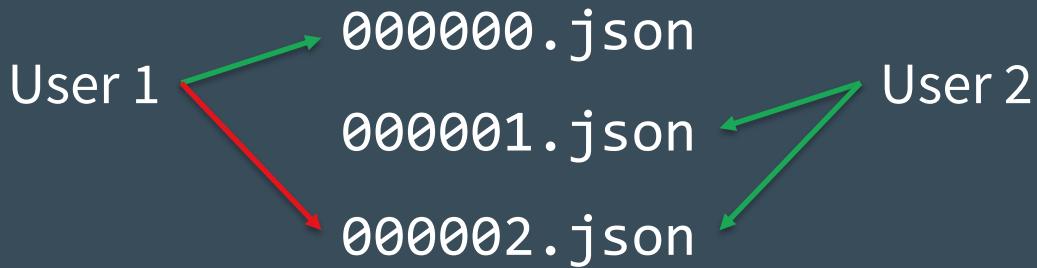
# Implementing Atomicity

Changes to the table  
are stored as  
*ordered, atomic*  
units called commits



# Ensuring Serializability

Need to agree on the order of changes, even when there are multiple writers.



# Solving Conflicts Optimistically

1. Record start version
2. Record reads/writes
3. Attempt commit
4. If someone else wins,  
check if anything you  
read has changed.
5. Try again.

Read: Schema

Write: Append

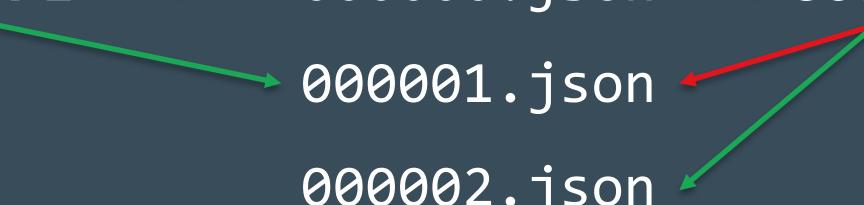
**User 1**



Read: Schema

Write: Append

**User 2**



# Handling Massive Metadata

Large tables can have millions of files in them! How do we scale the metadata? Use Spark for scaling!

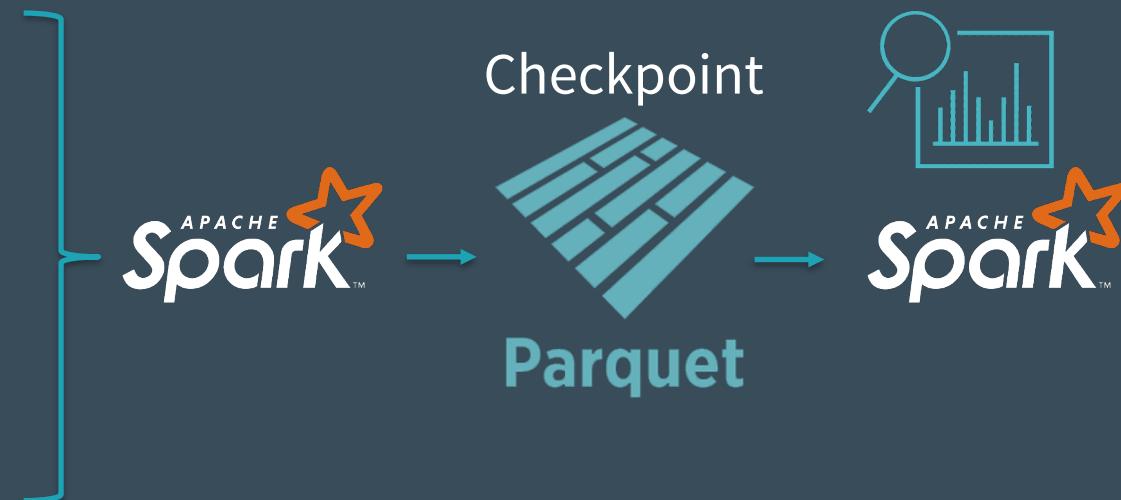
Add 1.parquet

Add 2.parquet

Remove 1.parquet

Remove 2.parquet

Add 3.parquet



Build your own Delta Lake  
at **<https://delta.io>**

